

# Estimate Parameters for Discrete Distributions

Fan Wang

2022-07-23

## Contents

<b>1</b>	<b>Binomial Probability Function Given Observed Discrete Probabilities</b>	<b>1</b>
1.1	Estimate Chance of Success given Fixed Number of Trials . . . . .	1
1.1.1	Approximating Transition Probability . . . . .	2
1.2	Estimate Chance of Success and the Number of Trials . . . . .	3
1.2.1	Optimal Theta Given M . . . . .	4
1.2.2	Optimal M given Optimal Theta Conditional on M . . . . .	10
1.2.3	Rescale Theta by Common Parameter . . . . .	14

## 1 Binomial Probability Function Given Observed Discrete Probabilities

Go to the [RMD](#), [R](#), [PDF](#), or [HTML](#) version of this file. Go back to [fan's REconTools](#) research support package, [R4Econ](#) examples page, [PkgTestR](#) packaging guide, or [Stat4Econ](#) course page.

### 1.1 Estimate Chance of Success given Fixed Number of Trials

We observe  $\{P(x_i)\}_{i=0}^N$ , where  $\sum_{i=0}^N P(x_i) = 1$ . Note this is observing probabilities not at  $N$  outcomes but  $N + 1$  outcomes starting at  $i = 0$ . What is the [binomial probability function](#) that would fit the observed data the best?

There are two parameters of the binomial distribution,  $N$ , which is the “number of trials”, and  $\theta$ , which is the probability of success for each trial. In this section, we treat  $N$  as already observed, so we only need to pick the  $\theta$  that provides the best fit between the binomial probabilities generated (given  $p$ ), and the data.

For the household that currently has zero child, one interpretation of this fitting procedure is that there are for this household  $N$  “trials”/chances of getting a child, and  $\theta$  is the chance of new arrival during each trial. By using the binomial distribution, we are assuming that the chance of new child arrival from across trials is independent, and the chances of arrival is identical during each trial. The interpretation is less clear when are looking at a household that already has more than zero children. Overall, the purpose here is to use a low-dimensional procedure to provide a good fit to the observed probabilities, and to be able to easily adjust the resulting parameters across observed discrete probabilities to change some aggregate outcome of interest (e.g., [fertility](#)).

We fit with a likelihood based approach. First, the likelihood function is:

$$\mathcal{L}(\theta) = \prod_{i=0}^N \left( \frac{N!}{(N-i)!i!} \cdot \theta^i \cdot (1-\theta)^{N-i} \right)^{P(x_i)}$$

Second, we write the log likelihood. Note that we observe  $P(x_i)$  from the data.

$$\begin{aligned}
\ln(\mathcal{L}(\theta)) &= \sum_{i=0}^N P(x_i) \ln \left( \frac{N!}{(N-i)!i!} \cdot \theta^i \cdot (1-\theta)^{N-i} \right) \\
&= \sum_{i=0}^N \left( P(x_i) \ln \left( \frac{N!}{(N-i)!i!} \right) + P(x_i)i \ln(\theta) + P(x_i)(N-i) \ln(1-\theta) \right)
\end{aligned}$$

Third, we take the derivative of the log likelihood with respect to the parameter of interest  $\theta$ .

$$\frac{d \ln(\mathcal{L}(\theta))}{d\theta} = \sum_{i=0}^N \left( \frac{P(x_i)i}{\theta} - \frac{P(x_i)(N-i)}{1-\theta} \right)$$

Fourth, we set the derivative above to be equal to zero.

$$\begin{aligned}
\frac{\sum_{i=0}^N P(x_i) \cdot i}{\theta} &= \frac{\sum_{i=0}^N P(x_i)(N-i)}{1-\theta} \\
\sum_{i=0}^N P(x_i) \cdot i - \theta \cdot \sum_{i=0}^N P(x_i) \cdot i &= \theta \cdot \sum_{i=0}^N P(x_i)(N-i) \\
\theta &= \frac{\sum_{i=0}^N P(x_i) \cdot i}{\sum_{i=0}^N (P(x_i)(N-i) + P(x_i) \cdot i)} \\
\theta &= \frac{\sum_{i=0}^N P(x_i) \cdot i}{N \cdot \sum_{i=0}^N P(x_i)} \\
\theta &= \frac{\sum_{i=0}^N P(x_i) \cdot i}{N}
\end{aligned}$$

Fifth, given the derivation above, we have found the maximizing parameter of the likelihood function. The maximizer is the ratio between the mean (probability-weighted average) of the observed discrete data, divided by the  $N$ , where  $N + 1$  is the number of discrete data points observed.

$$\arg \max_{\theta} \ln(\mathcal{L}(\theta)) = \frac{\sum_{i=0}^N P(x_i) \cdot i}{N}$$

Suppose that the observe probability data is measured without error from a binomial distribution, in that case  $\sum_{i=0}^N P(x_i = i) \cdot i = \mu_{\text{binomial}} = N \cdot \theta$ , which means the mean number of “wins” is the number of “trials” times the probability of winning each trial.

### 1.1.1 Approximating Transition Probability

As an example, we now proceed with an example. Below we have a 5 by 5 transition matrix, where each of the row sums up to 1. Fixing  $N = 4$ , we find the best fitting  $\theta$  parameters for each row of the transition matrix.

This matrix example below is used as the household structure kids transition matrix for a particular set of 18 year old individuals [Nygaard, Sørensen, and Wang \(2022\) \(preprint pdf\)](#). Each row corresponds to a different household size in the current period, and each column is the chance of transitioning to different household sizes in the next period.

First, generate the input matrixes.

```

# Construct transition matrix (row sums to 1, 5 by 5)
mt_pi_kids_trans <- matrix(
  data = c(
    0.8929, 0.1045, 0.0026, 0.0000, 0.0000,
    0.0547, 0.6704, 0.2708, 0.0040, 0.0001,
    0.0014, 0.0571, 0.7931, 0.1462, 0.0022,
    0.0000, 0.0009, 0.0748, 0.7885, 0.1358,
    0.0000, 0.0000, 0.0008, 0.0593, 0.9399
  ),
  nrow = 5, ncol = 5,
  byrow = TRUE
)
# check row sum
rowSums(mt_pi_kids_trans)

```

```
## [1] 1 1 1 1 1
```

```

# Construct the x vector, this is (5 by 1)
# we have x_0, to x_4, for the 5 columns
# x_0 = 0, and x_4 = 4
ar_x <- matrix(c(0, 1, 2, 3, 4), nrow = 5, ncol = 1)

```

Second, using our formula from above, estimate the  $\theta$  parameters for each household size (row) this period.

```

# (5 by 5) times (5 by 1) generates (5 by 1)
# they correspond to 5 different  $\sum_{i=0}^N P(x_i)x_i$ 
# Note divide by 4 not 5
ar_row_expected_x <- (mt_pi_kids_trans %*% ar_x) / (4)

# estimated thetas have been found
ar_theta_esti <- ar_row_expected_x
print(ar_theta_esti)

```

```

##           [,1]
## [1,] 0.027425
## [2,] 0.306100
## [3,] 0.522675
## [4,] 0.764800
## [5,] 0.984775

```

Third, generate the binomial probabilities and construct the 5 by 5 matrix based on that the row specific estimated  $\theta$  values and the shared and fixed value of  $N = 4$ .

```

# Generate Binomial
mt_dbinom_approx <- t(sapply(
  ar_theta_esti, dbinom,
  x = seq(0, 4), size = 4
))
# Present
print(round(mt_dbinom_approx, 3))

```

```

##           [,1] [,2] [,3] [,4] [,5]
## [1,] 0.895 0.101 0.004 0.000 0.000
## [2,] 0.232 0.409 0.271 0.080 0.009
## [3,] 0.052 0.227 0.373 0.273 0.075
## [4,] 0.003 0.040 0.194 0.421 0.342

```

## [5,] 0.000 0.000 0.001 0.058 0.940

Evaluating the results, we see that this provides a good fit for the first and fifth rows, but the fit for the three middle rows seem to be substantially worse. The observed transition mass is much more persistent than the fitted transition.

## 1.2 Estimate Chance of Success and the Number of Trials

In contrast to the prior section, in this section, we will allow for both the “number of trials” and the “chance of success” parameters to vary at the same time to provide a better fit between predicted and observed transition probabilities.

Since the Binomial distribution is approximately normal, so in principle, we should be able to use the binomial to fit a discretized normal random variables—if each row of the transition matrix is discretized from some underlying normal process. A problem with the prior section is that the “number of trials” was fixed at  $N = 4$ , which is very small. At this small  $N$ , the binomial and normal approximation does not work well.

We assume that the observed probabilities are aggregated based on underlying neighboring subgroups. Suppose there is an underlying discrete random variable with 20 incremental subgroups from 0 to 19, perhaps we observe an aggregated version of these where probabilities for each four consecutive subgroups is summed up, forming five aggregated groups. Since we only observe the five aggregated probabilities, we observe a “fuzzy” version of true underlying data with 20 subgroup probabilities.

To be consistent with the prior example,  $N + 1$  is still the number of discrete observed probability categories. New parameter  $M \in \mathbb{N}$  is an integer that determines how many discrete evenly-spaced sub-categories are within each of the observed  $N + 1$  probability categories. This means the binomial “number of trials” is equal to

$$\text{numberOfBinomialTrials}(N, M) = (N + 1) \cdot M - 1$$

To allow for changing the “number of trials” parameter, we adjust our likelihood-based estimation objective function. Given  $N = 4$ , rather than fixing the “number of trials” at 4, we now allow it to be 9 (if  $M = 2$ ), 14 (if  $M = 3$ ), 19 (if  $M = 4$ ),  $\dots$ , and see how well the sum of probabilities from neighboring subgroups match with the observed probabilities. Our new likelihood function is:

$$\mathcal{L}(\theta, M) = \prod_{i=0}^N \left( \sum_{j=0}^{M-1} P_{\text{binom}} \left( \underbrace{(i \times M + j)}_{\text{Number of success}} ; \overbrace{\theta}^{\text{succ. prb.}}, \underbrace{(N + 1) \times M - 1}_{\text{Number of trials}} \right) \right)^{P(x_i)}$$

In the equation above,  $P_{\text{binom}}$  follows the same binomial probability formula as before, it is evaluated at different “number of successes”, given parameter  $\theta$  (success probability) and parameter  $M$ , which determines the “number of trials” given  $N$ .

Given one set of observed probabilities to fit, the constrained maximum likelihood optimization problem is now:

$$\max_{\substack{\theta \in (0,1) \\ M \in \mathbb{N}}} \mathcal{L}(\theta, M)$$

When we have  $K$  sets of observed probabilities as in the example earlier with the five transition probabilities, we could allow for  $K$  different  $\theta$  parameters, but constrain all  $K$  transitions to share the same  $M$ . In that case, the constrained optimization problem becomes:

$$\max_{\substack{\{\theta_k \in (0,1)\}_{k=1}^K \\ M \in \mathbb{N}}} \prod_{k=1}^K \mathcal{L}(\theta_k, M)$$

We do not, however, have analytical solutions for this optimization problem. We will solve for this problem numerically.

### 1.2.1 Optimal Theta Given M

For the numerical procedure, given  $M$ , we will find the best fitting  $\theta$ , then we increment  $M$ , and we compare fit and best fitting  $\theta$  across  $M$ .

First, we construct the likelihood function that considers the sum of probability between non-overlapping  $M$  adjacent groups. The function computes the log likelihood, given  $N + 1$  points of observed probabilities,  $\theta$  values, and  $M$ .

```
#' Probability group aggregator
#' @param it_n Integer such that it_n + 1 is the number of observed
#' discrete outcomes.
#' @param it_m Integer such that \eqn{(it_n + 1) * it_m - 1} is the
#' number of trials for a binomial experiment.
ffi_pbinom_group_theta_m <-
  function(fl_chance_success = 0.52,
           it_n = 4,
           it_m = 8) {

    # nbrtrl = number of trials
    it_nbrtrl <- (it_n + 1) * it_m - 1

    # Evaluate probability
    ar_dbinom <- dbinom(
      seq(0, it_nbrtrl),
      size = it_nbrtrl, prob = fl_chance_success
    )
    mt_dbinom <- matrix(
      ar_dbinom,
      nrow = it_n + 1, ncol = it_m,
      byrow = TRUE
    )

    # Sum probabilities each sub-groups
    ar_dbinom_grouped <- rowSums(mt_dbinom)

    return(ar_dbinom_grouped)
  }

# log-likelihood generator
# at the default value, the generated probabilities are similar to observed.
ffi_pbinom_lnlk_theta_m <-
  function(ar_drm_prob = c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022),
           fl_chance_success = 0.50,
           it_m = 1) {

    # it_n = (N+1)*M-1
    it_n <- length(ar_drm_prob) - 1

    # Call the grouping function
    ar_dbinom_grouped <- ffi_pbinom_group_theta_m(
      fl_chance_success = fl_chance_success,
      it_n = it_n,
```

```

    it_m = it_m
  )

  # Log of group-probabilities
  ar_dbinom_grouped_ln <- log(ar_dbinom_grouped)

  # Log likelihood
  fl_ln_likelihood <- ar_dbinom_grouped_ln %*% ar_drm_prob
  return(fl_ln_likelihood)
}

```

Second, test the function above with different theta values. Note that given the input example probability array and fixing  $M = 1$ , choosing among 9 different  $\theta$  values from 0.1 to 0.9, we find that  $\theta = 0.5$  provides the highest likelihood. It is clear that the likelihood function is single-peaked and smooth in the example here as we vary over  $\theta$  here.

```

# Test the function with different theta values
ar_drm_prob <- c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022)
ar_fl_theta <- seq(0.1, 0.9, length.out = 9)
it_m <- 1
ar_ln_likelihood <- sapply(
  ar_fl_theta,
  ffi_pbinom_lnlk_theta_m,
  ar_drm_prob = ar_drm_prob,
  it_m = it_m
)

# Label and print
mt_theta_likelihood <- cbind(ar_fl_theta, ar_ln_likelihood)
colnames(mt_theta_likelihood) <- c("theta", "log_likelihood")
kable(as_tibble(mt_theta_likelihood) %>%
  mutate(rank = min_rank(desc(log_likelihood))), caption = paste(
    "Log likelihood evaluated at varying",
    "theta values, given some vector",
    "of observed probabilities, and",
    paste0("M=", it_m),
    separator = " "
  )) %>% kable_styling_fc()

```

Log likelihood evaluated at varying theta values, given some vector of observed probabilities, and  $M=1$

theta	log_likelihood	rank
0.1	-3.312301	9
0.2	-2.088022	7
0.3	-1.495267	5
0.4	-1.188130	3
0.5	-1.069711	1
0.6	-1.114579	2
0.7	-1.341568	4
0.8	-1.836548	6
0.9	-2.913725	8

Third, construct likelihood with different values of  $M$ . Note that we are using  $\theta = 0.5$ , which we found from the previous step as the likelihood maximizing parameter value given  $M = 1$ . The table that we generate below shows that the likelihood is maximized at  $M = 7$ . It is clear that the likelihood function is single-peaked

and smooth in the example here as we vary over  $M$ .

```
# Test the function with different theta values
ar_drm_prob <- c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022)
fl_theta <- 0.50
ar_it_m <- seq(1, 10)
ar_ln_likelihood <- sapply(
  ar_it_m,
  ffi_pbinom_lnlk_theta_m,
  ar_drm_prob = ar_drm_prob,
  fl_chance_success = fl_theta
)
# Label and print
mt_m_likelihood <- cbind(ar_it_m, ar_ln_likelihood)
colnames(mt_m_likelihood) <- c("M", "log_likelihood")
kable(as_tibble(mt_m_likelihood) %>%
  mutate(rank = min_rank(desc(log_likelihood))), caption = paste(
    "Log likelihood evaluated at varying",
    "M values, given some vector",
    "of observed probabilities, and",
    paste0("theta=", fl_theta),
    separator = " "
  )) %>% kable_styling_fc()
```

Log likelihood evaluated at varying M values, given some vector of observed probabilities, and theta=0.5

M	log_likelihood	rank
1	-1.0697106	10
2	-0.8713480	9
3	-0.7772675	8
4	-0.7266149	7
5	-0.6988888	6
6	-0.6850268	3
7	-0.6803342	1
8	-0.6820966	2
9	-0.6886103	4
10	-0.6987371	5

Fourth, we develop a grid evaluator of the function we just created. Given a vector of different  $\theta$  values, we find the index that maximizes the likelihood, and we create new lower and upper bound that are around the maximizing grid point for  $\theta$ .

```
# Evaluate likelihood along a grid and find the
# surrounding points around the maximizing grid point.
ffi_pbinom_lnlk_theta_m_grid <-
  function(ar_drm_prob = c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022),
    it_m = 8,
    it_p_success_grid_len = 10,
    fl_min_p_success = 1e-5,
    fl_max_p_success = 1 - 1e-5) {

    # Theta array rebuilt
    ar_fl_theta <- seq(
      fl_min_p_success,
      fl_max_p_success,
```

```

        length.out = it_p_success_grid_len
    )

    # Evaluate likelihood
    ar_ln_likelihood <- sapply(
        ar_fl_theta,
        ffi_pbinom_lnlk_theta_m,
        ar_drm_prob = ar_drm_prob,
        it_m = it_m
    )

    # Find min grid
    it_max_idx <- which.max(ar_ln_likelihood)
    fl_max_val <- ar_ln_likelihood[it_max_idx]

    # Find lower and upper bound
    fl_min_p_success_new <- ar_fl_theta[
        max(it_max_idx - 1, 1)
    ]
    fl_max_p_success_new <- ar_fl_theta[
        min(it_max_idx + 1, it_p_success_grid_len)
    ]

    # return
    return(list(
        fl_max_val = fl_max_val,
        fl_min_p_success_new = fl_min_p_success_new,
        fl_max_p_success_new = fl_max_p_success_new
    ))
}

# test
ffi_pbinom_lnlk_theta_m_grid()

```

```

## $fl_max_val
## [1] -0.7219323
##
## $fl_min_p_success_new
## [1] 0.4444456
##
## $fl_max_p_success_new
## [1] 0.6666633

```

Fifth, given  $N + 1$  points of data, and given  $M$ , we find the best fit  $\theta$  value. We loop iteratively over successively finer and zoomed-in grids several times. For an example of this iterative-grid algorithm, see [Find Maximum By Iterating Over Grids](#).

```

# Find likelihood maximizing theta, given M
ffi_pbinom_lnlk_estitheta_fixm <-
    function(ar_drm_prob = c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022),
        it_m = 1,
        verbose = FALSE) {

    # Initialize min and max and tolerance criteria
    fl_min_p_success_cur <- 1e-5
    fl_max_p_success_cur <- 1 - 1e-5
    it_p_success_grid_len <- 10

```



```

fl_tol <- 1e-6
it_max_iter <- 10

# Initialize initial gaps etc
fl_gap <- 1e5
fl_min_ln_like_last <- 1e5
it_iter <- 0

# Iteratively loop over grid to find the maximum by zooming in
while ((fl_gap > fl_tol) && it_iter <= it_max_iter) {

  # Iterator counts up
  it_iter <- it_iter + 1
  if (verbose) print(paste0("it_iter=", it_iter))

  # build array
  ls_find_max <- ffi_pbinom_lnlk_theta_m_grid(
    ar_drm_prob = ar_drm_prob,
    it_m = it_m,
    it_p_success_grid_len = it_p_success_grid_len,
    fl_min_p_success = fl_min_p_success_cur,
    fl_max_p_success = fl_max_p_success_cur
  )

  # Min objective value current
  fl_max_ln_like <- ls_find_max$fl_max_val
  # Find new lower and upper bound
  fl_min_p_success_cur <- ls_find_max$fl_min_p_success_new
  fl_max_p_success_cur <- ls_find_max$fl_max_p_success_new
  if (verbose) print(paste0("min_p_succ_cur=", fl_min_p_success_cur))
  if (verbose) print(paste0("max_p_succ_cur=", fl_max_p_success_cur))

  # Compare
  fl_gap <- abs(fl_max_ln_like - fl_min_ln_like_last)
  fl_max_ln_like_last <- fl_max_ln_like
  if (verbose) print(paste0("fl_gap=", fl_gap))
}

# Find estimate for theta
fl_p_success_argmax <- (fl_min_p_success_cur + fl_max_p_success_cur) / 2

# return
return(list(
  fl_p_success_argmax = fl_p_success_argmax,
  fl_max_ln_like = fl_max_ln_like
))
}

# Test with different M values
ar_drm_prob <- c(0.0014, 0.0571, 0.7931, 0.1462, 0.0022)
ffi_pbinom_lnlk_estitheta_fixm(ar_drm_prob = ar_drm_prob, it_m = 1)

## $fl_p_success_argmax
## [1] 0.522675
##

```

```
## $fl_max_ln_like
## [1] -1.065596
ffi_pbinom_lnlk_estitheta_fixm(ar_drm_prob = ar_drm_prob, it_m = 5)

## $fl_p_success_argmax
## [1] 0.5195473
##
## $fl_max_ln_like
## [1] -0.685463
ffi_pbinom_lnlk_estitheta_fixm(ar_drm_prob = ar_drm_prob, it_m = 9)

## $fl_p_success_argmax
## [1] 0.5211407
##
## $fl_max_ln_like
## [1] -0.665317
```

### 1.2.2 Optimal M given Optimal Theta Conditional on M

We now proceed to find the optimum  $M$  value, assuming that multiple probability vectors can have different  $\theta$  but must share  $M$ .

First, we write the function that finds the optimizing  $\theta$  that differs for each of the discrete probability vectors.

```
# Find likelihood maximizing theta, given M
ffi_pbinom_lnlk_estijnt <-
  function(mt_drm_prob = matrix(
    data = c(
      0.8929, 0.1045, 0.0026, 0.0000, 0.0000,
      0.0547, 0.6704, 0.2708, 0.0040, 0.0001
    ),
    nrow = 2, ncol = 5, byrow = TRUE
  ),
  it_m_max = 10,
  verbose = FALSE) {

  # Initialize min and max and tolerance criteria
  it_drm_set <- dim(mt_drm_prob)[1]
  ar_it_m <- seq(1, it_m_max)

  for (it_i in seq(1, it_drm_set)) {
    ls_estitheta_fixm_res <- sapply(
      ar_it_m,
      ffi_pbinom_lnlk_estitheta_fixm,
      ar_drm_prob = mt_drm_prob[it_i, ]
    )

    mt_estitheta_fixm_res <- t(matrix(
      as.numeric(ls_estitheta_fixm_res),
      nrow = 2, ncol = length(ls_estitheta_fixm_res) / 2
    ))
    mt_estitheta_fixm_res <- cbind(mt_estitheta_fixm_res, ar_it_m)
    colnames(mt_estitheta_fixm_res) <- c("esti_theta", "ln_like", "M")

    tb_estitheta_fixm_res <-
```

```

      as_tibble(mt_estitheta_fixm_res) %>%
      mutate(drm_group = it_i)

    if (it_i > 1) {
      tb_estitheta_fixm_res_stack <-
        bind_rows(
          tb_estitheta_fixm_res_stack,
          tb_estitheta_fixm_res
        )
    } else {
      tb_estitheta_fixm_res_stack <- tb_estitheta_fixm_res
    }
  }

  # Group max
  tb_estitheta_fixm_res_stack <- tb_estitheta_fixm_res_stack %>%
    group_by(drm_group) %>%
    mutate(drm_group_rank = min_rank(desc(ln_like))) %>%
    ungroup()

  # Overall M that maximizes sum of log likelihood
  tb_estitheta_fixm_res_stack <- tb_estitheta_fixm_res_stack %>%
    group_by(M) %>%
    mutate(ln_like_M_sum = sum(ln_like)) %>%
    ungroup() %>%
    mutate(drm_jnt_rank = dense_rank(desc(ln_like_M_sum))) %>%
    arrange(drm_group, M)

  # Return
  return(tb_estitheta_fixm_res_stack)
}

```

Second, we test the function above. In the table show below, we look at the *drm* (discrete random variable) coming from two rows of the transition matrix. For each row, we estimate the optimal  $\theta$  given different  $M$  values. By default, we try  $M$  from 1 through 10, the maximum value of  $M$  is controlled by the *it\_m\_max* parameter. We have for each *drm* and for each  $M$  the optimal  $\theta$  estimate along with the likelihood achieved. Then we find, for each *drm*, the likelihood maximizing  $M$  value. In this exercise, we are not constraining the two transition matrix rows to share the same  $M$ .

We find that for the first row of the transition matrix (*drm-group*=1),  $M = 2$  along with  $\theta = 0.063638$  achieves the highest likelihood. Likelihoods do not vary substantially across  $M$  for the first row. For the second row of the transition matrix (*drm-group*=2),  $M = 5$  along with  $\theta = 0.338998$  achieves the highest likelihood. Likelihoods vary more across  $M$  for this second row.

```

# Test with first two rows of transition matrix
tb_rest_res <- ffi_pbinom_lnlk_estijnt(mt_drm_prob = mt_pi_kids_trans[1:2, ])
kable(tb_rest_res %>% select(-ln_like_M_sum, -drm_jnt_rank),
      caption = paste(
        "Given 2 sets of discrete probabilities (drm-group=1 and =2),",
        "for each drm-group and for each M,",
        "show the estimated theta (esti-theta) that maximizes likelihood,",
        "as well as the maximum likelihood achieved (ln-like).",
        "Show within drm-group likelihood desc rank (drm-group-rank);",
        "There is a (drm-group-rank=1) for each (drm-group),",
        "where we maximize likelihood by adjusting both M and theta",

```

```

    "at the same time.",
    "M rows for each set of (drm-group).",
    "drm = discrete random variable.",
    separator = " "
  )
) %>% kable_styling_fc()

```

Given 2 sets of discrete probabilities (drm-group=1 and =2), for each drm-group and for each M, show the estimated theta (esti-theta) that maximizes likelihood, as well as the maximum likelihood achieved (ln-like). Show within drm-group likelihood desc rank (drm-group-rank); There is a (drm-group-rank=1) for each (drm-group), where we maximize likelihood by adjusting both M and theta at the same time. M rows for each set of (drm-group). drm = discrete random variable.

esti_theta	ln_like	M	drm_group	drm_group_rank
0.0274250	-0.3531696	1	1	2
0.0636388	-0.3529403	2	1	1
0.0847539	-0.3545155	3	1	3
0.0984937	-0.3565668	4	1	4
0.1082890	-0.3587232	5	1	5
0.1157172	-0.3608758	6	1	6
0.1215997	-0.3629937	7	1	7
0.1264084	-0.3650707	8	1	8
0.1304352	-0.3671085	9	1	9
0.1338718	-0.3691108	10	1	10
0.3061000	-1.0436529	1	2	10
0.3264671	-0.8935429	2	2	9
0.3331335	-0.8329513	3	2	5
0.3366341	-0.8092757	4	2	2
0.3389987	-0.8042176	5	2	1
0.3408714	-0.8096444	6	2	3
0.3425018	-0.8213990	7	2	4
0.3439940	-0.8371529	8	2	6
0.3453910	-0.8555141	9	2	7
0.3467095	-0.8756103	10	2	8

Third, we now look at all five rows of the transition matrix together (five different drm-groups). For each M, we add up the likelihood from different drm-groups (shown in ln-like-M-sum), and we find the M that provides that highest overall likelihood. In this case, M = 6 achieves the highest likelihood, with estimated  $\theta$  values at 0.1157172, 0.3408714, 0.5197787, 0.7214638, and 0.9009770 for each one of the transition probability rows from the first to the last.

```

# Test with all rows of transition matrix
tb_rest_res <- ffi_pbinom_lnlk_estijnt(mt_drm_prob = mt_pi_kids_trans)
kable(tb_rest_res %>% filter(drm_jnt_rank == 1),
  caption = paste(
    "Given 2 sets of discrete probabilities (drm-group=1 and =2).",
    "pick the M",
    "where the overall likelihood (ln-like-M-sum) summed across",
    "likelihoods (ln-like) from the two sets of discrete probabilities",
    "is maximized (drm-jnt-rank == 1).",
    "Theta differs for each drm-group, but M is shared.",
    "drm = discrete random variable.",
    separator = " "
  )
)

```

```
)
) %>% kable_styling_fc()
```

Given 2 sets of discrete probabilities (drm-group=1 and =2), pick the M where the overall likelihood (ln-like-M-sum) summed across likelihoods (ln-like) from the two sets of discrete probabilities is maximized (drm-jnt-rank == 1). Theta differs for each drm-group, but M is shared. drm = discrete random variable.

esti_theta	ln_like	M	drm_group	drm_group_rank	ln_like_M_sum	drm_jnt_rank
0.1157172	-0.3608758	6	1	6	-2.734403	1
0.3408714	-0.8096444	6	2	3	-2.734403	1
0.5197787	-0.6693215	6	3	4	-2.734403	1
0.7214638	-0.6602480	6	4	1	-2.734403	1
0.9009770	-0.2343137	6	5	6	-2.734403	1

Fourth, we evaluate the predicted probabilities given our optimizing parameters and compare. Now the predictions and the actual results are much closer, compared to the results when we fixed the “number of trials” at 4 from the prior section. In this section here, we are using 6 parameters to fit 20 observed probabilities (excluding the last category). In the previous section we used 5 parameters to fit 20 observed probabilities. We achieved substantial gain in fit by adding in one more parameter. If we are willing to allow  $M$  to differ for each of the transition probability rows, we would be using 10 parameters to fit 20 observed probabilities and achieve an even closer fit.

```
# Estimate for full data input matrix (all 5 rows)
tb_rest_res <- ffi_pbinom_lnlk_estijnt(mt_drm_prob = mt_pi_kids_trans)

# Compute predicted probabilities at maximizing parameters
mt_dbinom_grouped <-
  apply(
    tb_rest_res %>% filter(drm_jnt_rank == 1),
    1,
    function(row) {
      ffi_pbinom_group_theta_m(
        fl_chance_success = row[["esti_theta"]],
        it_n = (dim(mt_pi_kids_trans)[1] - 1),
        it_m = row[["M"]]
      )
    }
  )

# Show predictions
kable(round(t(mt_dbinom_grouped), 3),
      caption = "Predicted transitions given estimates"
) %>% kable_styling_fc()
```

Predicted transitions given estimates

0.889	0.111	0.000	0.000	0.000
0.038	0.702	0.258	0.002	0.000
0.000	0.092	0.724	0.183	0.001
0.000	0.000	0.081	0.777	0.141
0.000	0.000	0.000	0.061	0.939

```
# Show data
kable(round(mt_pi_kids_trans, 3),
      caption = "Data transitions probabilities")
```

```
) %>% kable_styling_fc()
```

Data transitions probabilities

0.893	0.104	0.003	0.000	0.000
0.055	0.670	0.271	0.004	0.000
0.001	0.057	0.793	0.146	0.002
0.000	0.001	0.075	0.788	0.136
0.000	0.000	0.001	0.059	0.940

```
# Show data and prediction differences
kable(round(mt_pi_kids_trans - t(mt_dbinom_grouped), 3),
      caption = "Data transitions probabilities - predictions"
) %>% kable_styling_fc()
```

Data transitions probabilities - predictions

0.004	-0.007	0.003	0.000	0.000
0.017	-0.032	0.013	0.002	0.000
0.001	-0.035	0.069	-0.037	0.002
0.000	0.001	-0.007	0.011	-0.006
0.000	0.000	0.001	-0.002	0.001

### 1.2.3 Rescale Theta by Common Parameter

Following our discussions here, [Rescale a Fraction Between 0 and 1](#), we now proceed to test a procedure where we adjust a  $\lambda$  parameter, which will lead to a rescaling of each of the estimated  $\theta$  value. Specifically, we have the following rescaling function:

$$\hat{\theta}(\theta, \lambda) = \theta \cdot \left( \frac{\exp(\lambda)}{1 + \theta \cdot (\exp(\lambda) - 1)} \right)$$

First, we define the rescaling function, identical to the rescaling function from [Rescale a Fraction Between 0 and 1](#).

```
# Construct the formula
ffi_theta_lambda_0t1 <- function(theta, lambda) {
  if (is.finite(exp(lambda))) {
    theta * (exp(lambda) / (1 + theta * (exp(lambda) - 1)))
  } else {
    # If lambda is large, exp(lambda)=inf, ratio above becomes 1
    1
  }
}
```

Second, we create  $\hat{\theta}(\theta, \lambda)$ , with different scaling  $\lambda$  values, for the optimizing  $\theta$  for each of the drm-groups. We can see that when  $\lambda < 0$ ,  $0 < \hat{\theta} < \theta$ , when  $\lambda = 0$ ,  $\hat{\theta} = \theta$ , and when  $\lambda > 0$ ,  $1 > \hat{\theta} > \theta$ . Given our transformation  $\hat{\theta} \in (0, 1)$  is between 0 and 1 for all  $\lambda \in \mathbb{R}$ . See [Rescale a Fraction Between 0 and 1](#) for more details on this rescaling function.

```
# Estimation results
tb_rest_res_opti <- ffi_pbinom_lnlk_estijnt(mt_drm_prob = mt_pi_kids_trans) %>%
  filter(drm_jnt_rank == 1) %>%
  select(esti_theta, M, drm_group, ln_like, ln_like_M_sum)
```

```

# Generate theta_hat columns
tb_rest_res_opti_theta_hat <- tb_rest_res_opti %>%
  mutate(
    theta_hat_lambda_n2 = ffi_theta_lambda_0t1(
      theta = esti_theta, lambda = -2
    ),
    theta_hat_lambda_n1 = ffi_theta_lambda_0t1(
      theta = esti_theta, lambda = -1
    ),
    theta_hat_lambda_zr = ffi_theta_lambda_0t1(
      theta = esti_theta, lambda = 0
    ),
    theta_hat_lambda_p1 = ffi_theta_lambda_0t1(
      theta = esti_theta, lambda = +1
    ),
    theta_hat_lambda_p2 = ffi_theta_lambda_0t1(
      theta = esti_theta, lambda = +2
    )
  ) %>%
  select(drm_group, esti_theta, contains("theta_hat"))
# Show data and prediction differences
kable(tb_rest_res_opti_theta_hat,
  caption = paste(
    "Theta transform with different common lambda.",
    "n2 is lambda = -2,",
    "zr is lambda = 0 (theta-hat=theta)",
    "p2 is lambda = +2.",
    separator = " "
  )
) %>% kable_styling_fc_wide()

```

Theta transform with different common lambda. n2 is lambda = -2, zr is lambda = 0 (theta-hat=theta), p2 is lambda = +2.

drm_group	esti_theta	theta_hat_lambda_n2	theta_hat_lambda_n1	theta_hat_lambda_zr	theta_hat_lambda_p1	theta_hat_lambda_p2
1	0.1157172	0.0174018	0.0459296	0.1157172	0.2623814	0.4915939
2	0.3408714	0.0654112	0.1598408	0.3408714	0.5843330	0.7925864
3	0.5197787	0.1277675	0.2847860	0.5197787	0.7463343	0.8888607
4	0.7214638	0.2595582	0.4879358	0.7214638	0.8756356	0.9503454
5	0.9009770	0.5518449	0.7699675	0.9009770	0.9611390	0.9853438